

▼ Introduction

In this assignment, you are going to build a classifier for named entities from the Groningen Meaning Bank corpus. Named entity recognition (NER) takes noun phrases from a text and identifies whether they are persons, organizations, and so on. You will be using the Groningen Meaning Bank named entity corpus available on mltgpu at /scratch/lt2222-v21-resources/GMB_dataset.txt. In this version of the task, you will assume we know *that* something is a named entity, and instead use multi-class classification to identify its type. So you will be doing named entity classification but *not* recognition.

The data looks like this:

3996	182.0	Nicole	NNP	B-per	
3997	182.0	Ritchie	NNP	I-per	
3998	182.0	is	VBZ	O	
3999	182.0	pregnant		JJ	O
4000	182.0	.	.	O	
4001	183.0	Speaking		VBG	O
4002	183.0	to	TO	O	
4003	183.0	ABC	NNP	B-org	
4004	183.0	News	NNP	I-org	
4005	183.0	interviewer		NN	O
4006	183.0	Dianne	NNP	B-per	
4007	183.0	Sawyer	NNP	I-per	
4008	183.0	,	,	O	
4009	183.0	the	DT	O	
4010	183.0	25-year-old		JJ	O
4011	183.0	co-star	NN	O	
4012	183.0	of	IN	O	
4013	183.0	TV	NN	O	
4014	183.0	's	POS	O	
4015	183.0	The	DT	B-art	
4016	183.0	Simple	NNP	I-art	
4017	183.0	Life	NNP	I-art	
4018	183.0	said	VBD	O	
4019	183.0	she	PRP	O	
4020	183.0	is	VBZ	O	
4021	183.0	almost	RB	O	
4022	183.0	four	CD	O	
4023	183.0	months	NNS	O	
4024	183.0	along	IN	O	
4025	183.0	in	IN	O	

4026	183.0	her	PRP\$	O
4027	183.0	pregnancy	NN	O
4028	183.0	.	.	O

The first column is the line number. The second column is a sentence number (for some reason given as a float; ignore it). The third column is the word. The fourth column is a part of speech (POS) tag in Penn Treebank format. The last column contains the named entity annotation.

The annotation works like this. Every `o` just means that the row does not represent a named entity. `B-xyz` means the first word in a named entity with type `xyz`. `I-xyz` means the second and later words of an `xyz` entity, if there are any. That means that every time there's a `B` or an `I`, there's a named entity.

The entity types in the corpus are `art`, `eve`, `geo`, `gpe`, `nat`, `org`, `per`, and `tim`

Your task is the following.

1. To preprocess the text (lowercase and lemmatize; punctuation can be preserved as it gets its own rows).
2. To create instances from every from every identified named entity in the text with the type of the NE as the class, and a surrounding context of five words on either side as the features.
3. To generate vectors and split the instances into training and testing datasets at random.
4. To train a support vector machine (via `sklearn.svm.LinearSVC`) for classifying the NERs.
5. To evaluate the performance of the classifier.

You will do this by modifying a separate file containing functions that will be called from this notebook as a module. You can modify this notebook for testing purposes but please only submit the original. You will document everything in Markdown in `README.md` and submit a GitHub repository URL.

This assignment is due on **Tuesday, 2021 March 9 at 23:59**. It has **25 points** and **7 bonus points**.

```
!pip install lazypredict #try to use lazypredict classifier for the first time (optio
!pip install scikit-plot
```

```
import nltk
nltk.download('stopwords')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Unzipping corpora/wordnet.zip.
True
```

```
import importlib
import a2
```

```
import sys
import pandas as pd

import numpy as np
import scikitplot as skplt

from sklearn.svm import LinearSVC
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import plot_confusion_matrix, confusion_matrix

import nltk

import matplotlib.pyplot as plt

import lazypredict
from lazypredict.Supervised import LazyClassifier
```

▼ Part 1 - preprocessing (3 points)

See step 1 above. The data is coming to you as an unused file handle object. You can return the data in any indexable form you like. You can also choose to remove infrequent or uninformative words to reduce the size of the feature space. (Document this in README.md.)

```
gmbfile = open('GMB_dataset.txt', "r")
inputdata = a2.preprocess(gmbfile)
inputdata[20:40]
```

	Sentence #	Word	POS	Tag	Tag_prefix	Tag_entity	Word_seq
36	2.0	banner	NNS	O	O	None	8
39	2.0	slogan	NNS	O	O	None	9
41	2.0	""""	``	O	O	None	10
42	2.0	bush	NNP	B-per	B	per	11
43	2.0	number	NN	O	O	None	12
44	2.0	one	CD	O	O	None	13
45	2.0	terrorist	NN	O	O	None	14
46	2.0	""""	``	O	O	None	15

```
inputdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 43977 entries, 0 to 66160
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sentence #      43977 non-null  object
1   Word            43977 non-null  object
2   POS             43977 non-null  object
3   Tag             43977 non-null  object
4   Tag_prefix      43977 non-null  object
5   Tag_entity      9736 non-null   object
6   Word_seq        43977 non-null  int64
dtypes: int64(1), object(6)
memory usage: 2.7+ MB
```

▼ Part 2 - Creating instances (7 points)

Do step 2 above. You will create a collection of Instance objects. Remember to consider the case where the NE is at the beginning of a sentence or at the end, or close to either (you can create a special start token for that). You can also start counting from before the **B** end of the NE mention and after the last **I** of the NE mention. That means that the instances should include things before and after the named entity mention, but not the named entity text itself.

```
n=5
```

```
instances = a2.create_instances(inputdata,n=n,skip_ne=False)
instances[20:40]
```

	0	1	2	3	4	5	6	
geo	percent	world	smoker	thought	live	india	alone	S
geo	world	smoker	thought	live	china	alone	S5	S
gpe	S1	S2	S3	S4	S5	military	launched	offensive
geo	S5	pakistani	military	launched	offensive	hunt	taliban	insurge
org	military	launched	offensive	orakzai	hunt	insurgent	S5	S
nat	die	cancer	soon	greater	death	tuberculosis	malaria	combine
tim	may	rise	27	million	year	17	million	peop
org	force	tried	suppress	report	abuse	prisoner	S5	S
org	S3	S4	S5	document	released	civil	liberty	unic
org	S4	S5	document	released	american	liberty	union	tuesda
org	S5	document	released	american	civil	union	tuesday	s
org	document	released	american	civil	liberty	tuesday	say	sta
tim	released	american	civil	liberty	union	say	staff	membr
org	union	tuesday	say	staff	member	defense	intelligence	agenc
org	tuesday	say	staff	member	pentagon	intelligence	agency	d
org	say	staff	member	pentagon	defense	agency	dia	witnesse
org	staff	member	pentagon	defense	intelligence	dia	witnessed	sever
org	member	pentagon	defense	intelligence	agency	witnessed	several	incide
org	S5	document	also	included	complaint	personnel	monitored	speci

```

encode_strings = ["S{}".format(i) for i in range(1, n + 1)]
instances = instances.replace(to_replace=encode_strings,value="NA")
instances[20:40]

```

	0	1	2	3	4	5	6	
geo	percent	world	smoker	thought	live	india	alone	N
geo	world	smoker	thought	live	china	alone	NA	N
gpe	NA	NA	NA	NA	NA	military	launched	offensiv
geo	NA	pakistani	military	launched	offensive	hunt	taliban	insurge
org	military	launched	offensive	orakzai	hunt	insurgent	NA	N
nat	die	cancer	soon	greater	death	tuberculosis	malaria	combine
tim	may	rise	27	million	year	17	million	peop
org	force	tried	suppress	report	abuse	prisoner	NA	N
org	NA	NA	NA	document	released	civil	liberty	unic
org	NA	NA	document	released	american	liberty	union	tuesda
org	NA	document	released	american	civil	union	tuesday	sæ
org	document	released	american	civil	liberty	tuesday	say	sta

▼ Part 3 - Creating the table and splitting (10 points)

Here you're going to write the functions that create a data table with "document" vectors representing each instance and split the table into training and testing sets and random with an 80%/20% train/test split.

```

org      staff  member  pentagon  defense  intelligence      dia  witnessed  sever

importlib.reload(a2)

bigdf = a2.create_table(instances,method='tfidf')
bigdf[20:40]
```

	class	10	100	103	10th	11	110	119	11th	12	120	123	12th	13
20	geo	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
21	geo	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
22	gpe	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
23	geo	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
24	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
25	nat	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
26	tim	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
27	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
28	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
29	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
30	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
31	org	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00

```

train_X, train_y, test_X, test_y = a2.ttsplit(bigdf)

print(train_X.shape, train_y.shape, test_X.shape, test_y.shape)

(7184, 4085) (7184,) (1796, 4085) (1796,)

len(test_y) / (len(test_y) + len(train_y))

0.2

len(test_X) / (len(test_X) + len(train_X))

0.2

test_y[0]

'per'

```

▼ Part 4 - Training the model (0 points)

This part you won't do yourself.

```

model = LinearSVC()
model.fit(train_X, train_y)

```

```
train_predictions = model.predict(train_X)
```

```
test_predictions = model.predict(test_X)
```

```
train_predictions
```

```
array(['tim', 'geo', 'gpe', ..., 'org', 'per', 'org'], dtype=object)
```

```
train_y
```

```
array(['gpe', 'geo', 'gpe', ..., 'org', 'per', 'org'], dtype=object)
```

```
test_predictions
```

```
array(['per', 'per', 'per', ..., 'org', 'geo', 'per'], dtype=object)
```

```
test_y
```

```
array(['per', 'org', 'gpe', ..., 'org', 'geo', 'per'], dtype=object)
```

```
model = LazyClassifier(verbose=1,ignore_warnings=True, custom_metric=None)
```

```
model.fit(train_X,test_X, train_y,test_y)
```

3%		1/29 [01:09<32:25, 69.49s/it]{'Model': 'AdaBoostClassifier', 'A
7%		2/29 [03:01<42:31, 94.51s/it]{'Model': 'BaggingClassifier', 'A
10%		3/29 [03:04<22:49, 52.66s/it]{'Model': 'BernoulliNB', 'Accuracy
14%		4/29 [44:18<7:00:16, 1008.66s/it]{'Model': 'CalibratedClassifi
21%		6/29 [45:23<2:51:32, 447.50s/it]{'Model': 'DecisionTreeClassifi
24%		7/29 [45:25<1:50:43, 301.98s/it]{'Model': 'DummyClassifier', 'A
28%		8/29 [45:30<1:12:29, 207.14s/it]{'Model': 'ExtraTreeClassifier
31%		9/29 [48:12<1:04:24, 193.20s/it]{'Model': 'ExtraTreesClassifier
34%		10/29 [48:16<42:39, 134.70s/it] {'Model': 'GaussianNB', 'Accur
38%		11/29 [51:03<43:25, 144.73s/it]{'Model': 'KNeighborsClassifier
41%		12/29 [51:21<30:06, 106.24s/it]{'Model': 'LabelPropagation', 'A
45%		13/29 [51:40<21:16, 79.80s/it] {'Model': 'LabelSpreading', 'Ac
48%		14/29 [54:56<28:42, 114.80s/it]{'Model': 'LinearDiscriminantAn
52%		15/29 [1:04:29<58:58, 252.78s/it]{'Model': 'LinearSVC', 'Accur
55%		16/29 [1:04:57<40:09, 185.36s/it]{'Model': 'LogisticRegression'
59%		17/29 [1:05:01<26:06, 130.54s/it]{'Model': 'NearestCentroid',
66%		19/29 [1:05:55<13:20, 80.07s/it]{'Model': 'PassiveAggressiveCla
69%		20/29 [1:06:15<09:17, 61.89s/it]{'Model': 'Perceptron', 'Accur
72%		21/29 [1:06:58<07:29, 56.21s/it]{'Model': 'QuadraticDiscrimina
76%		22/29 [1:07:53<06:31, 55.91s/it]{'Model': 'RandomForestClassifi
79%		23/29 [1:08:08<04:21, 43.61s/it]{'Model': 'RidgeClassifier', 'A
83%		24/29 [1:11:39<07:49, 93.84s/it]{'Model': 'RidgeClassifierCV',
86%		25/29 [1:13:08<06:09, 92.44s/it]{'Model': 'SGDClassifier', 'Acc
90%		26/29 [1:31:52<20:05, 401.98s/it]{'Model': 'SVC', 'Accuracy':
97%		28/29 [1:49:25<07:39, 459.36s/it]{'Model': 'XGBClassifier', 'A
100%		29/29 [1:49:52<00:00, 227.33s/it]{'Model': 'LGBMClassifier', 'A

(Accuracy	...	Time Taken
Model			

ExtraTreesClassifier	0.55	...	162.54
RandomForestClassifier	0.52	...	55.21
RidgeClassifierCV	0.45	...	211.01
NearestCentroid	0.46	...	3.07
LinearDiscriminantAnalysis	0.43	...	195.70
LogisticRegression	0.46	...	28.77
RidgeClassifier	0.44	...	14.92
ExtraTreeClassifier	0.42	...	4.07
GaussianNB	0.36	...	3.71
Perceptron	0.44	...	19.54
BaggingClassifier	0.46	...	112.03
PassiveAggressiveClassifier	0.44	...	52.21
XGBClassifier	0.46	...	1052.57
DecisionTreeClassifier	0.42	...	62.49
LGBMClassifier	0.46	...	27.40
SGDClassifier	0.48	...	89.15
LinearSVC	0.39	...	572.54
KNeighborsClassifier	0.37	...	167.48
SVC	0.48	...	1124.14
BernoulliNB	0.49	...	2.86
CalibratedClassifierCV	0.44	...	2474.19
QuadraticDiscriminantAnalysis	0.13	...	42.95
LabelSpreading	0.01	...	18.95
LabelPropagation	0.01	...	18.19
AdaBoostClassifier	0.28	...	69.49
DummyClassifier	0.20	...	2.39

[26 rows x 5 columns],

	Accuracy	...	Time Taken
Model		...	


▼ Part 5 - Evaluation (5 points)

Investigate for yourself what a "confusion matrix". Then implement a function that takes the data and produces a confusion matrix in any readable form that allows us to compare the performance of the model by class.

```
model.score(train_X,train_y),model.score(test_X,test_y)

(0.9001948775055679, 0.5005567928730512)
```

```
#Method 1
importlib.reload(a2)
a2.confusion_matrix(test_y, test_predictions)
```



#Method 2

```
np.set_printoptions(precision=2)
```

```
fig = plt.figure(figsize=(8,8))
```


```
class_names = bigdf['_class_'].unique()
```

```
disp = plot_confusion_matrix(model,test_X, test_y,
                             display_labels=class_names,
                             normalize='true')
```

```
confusion_matrix(test_y,test_predictions)
```

```
array([[ 6,  0,  0,  0,  0,  0,  5,  1],
       [ 0,  6,  2,  0,  1,  3,  4,  2],
       [ 1,  5, 249, 32,  0, 42, 54, 66],
       [ 0,  1, 60, 63,  0, 34, 46, 29],
       [ 0,  0,  5,  1,  0,  2,  1,  0],
       [ 0,  0, 68, 23,  0, 210, 46, 28],
       [ 1,  2, 56, 29,  0, 49, 272, 22],
       [ 1,  2, 72, 19,  1, 41, 40, 93]])
```

<Figure size 576x576 with 0 Axes>




Examine the matrix and describe your observations in README.md. In particular, what do you notice about the predictions on the training data compared to those on the test data.

▼ Bonus Part A - Error analysis (2 points)

Look at the weakest-performing classes in the confusion matrix (or any, if they all perform poorly to the same extent). Find some examples in the test data on which the classifier classified incorrectly for those classes. What do you think is the reason why those are hard? Consider linguistic factors and statistical factors, if applicable. Write your answer in README.md.


```
inputdata['Tag_entity'].value_counts(normalize=True).plot.bar()  
plt.xlabel('NE class')  
plt.ylabel('Number of instances')
```

Text(0, 0.5, 'Number of instances')



```
counts = bigdf['_class_'].value_counts(normalize=True)*100  
counts.plot.bar()  
plt.xlabel('NE class')  
plt.ylabel('Percentage of instances')
```

Text(0, 0.5, 'Percentage of instances')



```
counts
```

```
geo    25.51
per    23.02
org    20.87
tim    15.24
gpe    13.32
eve     0.87
art     0.86
nat     0.31
Name: _class_, dtype: float64
```

▼ Bonus Part B - Expanding the feature space (7 points)

Run the entire process above, but incorporate part-of-speech tag information into the feature vectors. It's your choice as to how to do this, but document it in README.md. Your new process should run from the single call below:

```
a2.bonusb( '/scratch/lt2222-v21-resources/GMB_dataset.txt' )
```